

Code

**Issues** 3.1k

Pull requests 866

Actions

Projects 6

Security

Insights

[New issue](#)[Jump to bottom](#)

# Unnecessary allocations in maximum and mapreduce

## #28752

**Open**

legendre6891 opened this issue on 19 Aug 2018 · 1 comment

Labels

iteration

performance



legendre6891 commented on 19 Aug 2018

In the function `f!` below, the memory allocation scales with `n`. In comparison, `g!` allocates no memory. The two are functionally the same (I think?).

```
function f!(x,y,z)
    n = length(x)
    for i in 1:n
        z[i] = maximum(a*b - b*i for (a,b) in zip(x,y))
    end
end

function g!(x,y,z)
    n = length(x)
    for i in 1:n

        m = -Inf
        for (a,b) in zip(x,y)
            v = a * b - b * i
            if v > m
                m = v
            end
        end
        z[i] = m

    end
end
```

Here is the test

```
julia> n = 500; x = randn(n); y = randn(n); z = similar(x);

julia> @allocated f!(x,y,z)
32000
```

I posted to Slack about this and it was suggested that I post file an issue. One theory is that `zip` is allocating in every iteration. In any case, `mapreduce` exhibits the same allocation behavior:

```
julia> function f!(x,y,z)
    for (i, v) in enumerate(x)
        z[i] = mapreduce(a->(a[1]*a[2]), max, zip(x,y), init = 0.0)
    end
end
f! (generic function with 1 method)
```

```
julia> @benchmark f!($a, $b, $c)
BenchmarkTools.Trial:
 memory estimate: 31.25 KiB
 allocs estimate: 1000
-----
 minimum time:      2.480 ms (0.00% GC)
 median time:      2.785 ms (0.00% GC)
 mean time:        2.922 ms (0.99% GC)
 maximum time:     39.308 ms (92.63% GC)
-----
 samples:          1708
 evals/sample:     1
```

```
julia> versioninfo(verbose=true)
Julia Version 1.0.0
Commit 5d4eaca0c9 (2018-08-08 20:58 UTC)
Platform Info:
 OS: Linux (x86_64-pc-linux-gnu)
 uname: Linux 4.14.61 #2 SMP Mon Aug 6 23:42:49 CDT 2018 x86_64 Intel(R) Core(TM) i7-8700K
CPU @ 3.70GHz
CPU: Intel(R) Core(TM) i7-8700K CPU @ 3.70GHz:
      speed        user        nice        sys        idle        irq
   #1-12  4700 MHz    779941 s    134 s    273512 s    29072680 s    0 s

Memory: 62.75409698486328 GB (60294.4453125 MB free)
Uptime: 25193.0 sec
Load Avg: 0.57763671875 0.50048828125 0.48681640625
WORD_SIZE: 64
LIBM: libopenlibm
LLVM: libLLVM-6.0.0 (ORCJIT, skylake)
```



**legendre6891** changed the title ~~Unnecessary allocations maximum and mapreduce~~ **Unnecessary allocations in maximum and mapreduce** on 19 Aug 2018

Some insightful experiments.

```
julia> versioninfo()
Julia Version 1.0.0
Commit 5d4eaca0c9 (2018-08-08 20:58 UTC)
Platform Info:
  OS: Windows (x86_64-w64-mingw32)
  CPU: Intel(R) Core(TM) i7-7700HQ CPU @ 2.80GHz
  WORD_SIZE: 64
  LIBM: libopenlibm
  LLVM: libLLVM-6.0.0 (ORCJIT, skylake)

julia> using BenchmarkTools

julia> a = rand(1000); b = rand(1000); c = rand(1000);
```

1. Firstly with the `zip` outside.

```
julia> function f!(x,y,z)
    n = length(x)
    xy = zip(x, y)
    for i in 1:n
        z[i] = maximum(a->(a[1]*a[2] - a[2]*i), xy)
    end
end

f! (generic function with 1 method)

julia> @benchmark f!($a, $b, $c)
BenchmarkTools.Trial:
  memory estimate: 32 bytes
  allocs estimate: 1
  -----
  minimum time:      2.459 ms (0.00% GC)
  median time:       2.672 ms (0.00% GC)
  mean time:         2.752 ms (0.00% GC)
  maximum time:      5.569 ms (0.00% GC)
  -----
  samples:           1815
  evals/sample:      1

julia> function g!(x,y,z)
    n = length(x)
    xy = zip(x, y)
    for i in 1:n
        z[i] = maximum((a[1]*a[2] - a[2] * i) for a in xy)
    end
end

g! (generic function with 1 method)

julia> @benchmark g!($a, $b, $c)
BenchmarkTools.Trial:
  memory estimate: 31.28 KiB
  allocs estimate: 1001
```

```

mean time:      2.337 ms (0.92% GC)
maximum time:   38.495 ms (91.85% GC)
-----
samples:        1701
evals/sample:   1

julia> function h!(x,y,z)
    n = length(x)
    xy = zip(x, y)
    for i in 1:n
        z[i] = mapreduce(a->(a[1]*a[2] - a[2]*i), max, xy, init=0.0)
    end
end

h! (generic function with 1 method)

julia> @benchmark h!($a, $b, $c)
BenchmarkTools.Trial:
  memory estimate: 32 bytes
  allocs estimate: 1
  -----
  minimum time:    2.441 ms (0.00% GC)
  median time:     2.711 ms (0.00% GC)
  mean time:       2.900 ms (0.00% GC)
  maximum time:    7.452 ms (0.00% GC)
  -----
  samples:         1721
  evals/sample:    1

```

As you can see, `maximum(i for i in I)` allocates in every iteration when semantically it shouldn't.

2. Further putting `zip` inside the loop causes even more allocations.

```

julia> function p!(x,y,z)
    n = length(x)
    for i in 1:n
        z[i] = mapreduce(a->(a[1]*a[2] - a[2]*i), max, zip(x, y), init=0.0)
    end
end

p! (generic function with 1 method)

julia> @benchmark p!($a, $b, $c)
BenchmarkTools.Trial:
  memory estimate: 31.25 KiB
  allocs estimate: 1000
  -----
  minimum time:    2.485 ms (0.00% GC)
  median time:     2.721 ms (0.00% GC)
  mean time:       2.818 ms (0.90% GC)
  maximum time:    36.950 ms (92.30% GC)
  -----
  samples:         1773
  evals/sample:    1

julia> function m!(x,y,z)
    n = length(x)
    for i in 1:n

```

```
n! (generic function with 1 method)
```

```
julia> @benchmark n!($a, $b, $c)
BenchmarkTools.Trial:
  memory estimate: 31.25 KiB
  allocs estimate: 1000
  -----
  minimum time:      2.421 ms (0.00% GC)
  median time:       2.669 ms (0.00% GC)
  mean time:         2.807 ms (0.96% GC)
  maximum time:      38.248 ms (93.12% GC)
  -----
  samples:           1780
  evals/sample:      1
```



```
julia> function n!(x,y,z)
    n = length(x)
    for i in 1:n
        z[i] = maximum((a[1]*a[2] - a[2]*i) for a in zip(x, y))
    end
end

n! (generic function with 1 method)
```

```
julia> @benchmark n!($a, $b, $c)
BenchmarkTools.Trial:
  memory estimate: 62.50 KiB
  allocs estimate: 2000
  -----
  minimum time:      2.564 ms (0.00% GC)
  median time:       2.848 ms (0.00% GC)
  mean time:         2.951 ms (1.03% GC)
  maximum time:      41.008 ms (93.47% GC)
  -----
  samples:           1693
  evals/sample:      1
```

 2



  **StefanKarpinski** added **performance** **iteration** labels on 21 Aug 2018

Assignees

No one assigned

Labels

**iteration** **performance**

Projects

None yet

---

No milestone

---

Linked pull requests

Successfully merging a pull request may close this issue.

None yet

---

3 participants

